

Suhas Hebbur Eshwar

(858) 250-9028 | shebbureshwar@ucsd.edu | linkedin.com/in/suhas-h-e/

EXPERIENCE

Oracle

Redwood City, CA

Senior Member of Technical Staff

Sept 2025 – Present

- Engineered a high-throughput index maintenance pipeline for external Apache Iceberg tables integrated with the Oracle Database engine, ensuring real-time data consistency and low-latency retrieval.
- Developed an incremental change-detection engine that processes file-level storage snapshot updates, eliminating the resource overhead of full index rebuilds.
- Optimized background maintenance scheduling across multi-instance clustered database deployments, enabling distributed workers to execute updates concurrently with minimal coordination overhead.
- Implemented a distributed task-queue locking mechanism within the database tier to provide synchronized access to pending work blocks across concurrent executors.
- Enhanced large-scale database generation and assignment performance by 30% by designing an in-memory caching layer that minimized repeated storage lookups during distributed partition assignments.

Member of Technical Staff

May 2024 – Aug 2025

- Implemented Float16/BFloat16 datatype support within the core Oracle database engine to accelerate high-performance vector search workloads.
- Engineered scalable storage APIs and custom optimization workflows that reduced the core engine's data storage footprint by over 50%.
- Developed metadata caching strategies, execution workflows, and query rewrite paths to support efficient relational system optimization and search execution.
- Optimized core database API ingestion layers using profiling and performance analysis tools, increasing write throughput by 47% for large-scale datasets of over 1M+ data rows.

Microsoft

Hyderabad, India

Software Engineer

July 2020 – Aug 2022

- Architected and engineered a global enterprise recruitment orchestration platform processing over 200k candidates monthly, designing customized multi-stage progression engines and automated workflows that reduced operational time-to-hire by 45%.
- Designed and implemented event-driven email and notification microservices using Apache Kafka and serverless Azure Functions, incorporating idempotent message processing and dead-letter queues (DLQs) to guarantee zero-loss, duplicate-free delivery.
- Managed global infrastructure deployments on Microsoft Azure, leveraging database read-replicas to optimize query performance while establishing automated backup strategies to maintain high availability.
- Steered core backend initiatives as technical lead, directly managing sprint allocations for junior engineers while establishing comprehensive testing frameworks and rigorous code quality standards.

EDUCATION

University of California, San Diego

San Diego, CA

Master of Science in Computer Science; GPA: 4.0/4.0

Sept 2022 – Mar 2024

- Relevant Coursework:** Distributed Systems, Recommender Systems, Data Mining, Big Data

PROJECTS

DropBox clone (UCSD) | Golang, Consistent Hashing, RAFT, gRPC

- Developed a cloud-based distributed file storage service with fault tolerance and synchronization capabilities.
- Implemented client-server interaction via gRPC and ensured data consistency across nodes using RAFT consensus.

TECHNICAL SKILLS

Languages: Python 3, C, Golang, SQL, C#/.NET Core, HTML/CSS, Java

Frameworks/Tools: Kafka, Docker, Azure, gRPC, Codex, Claude Code, ReactJS, Angular, Flask

Specializations: Distributed Systems, Database Internals, High Performance Systems, Vector Search, Full Stack Development